AI Opportunity Report

Industry context, company alignment and recommended use case

Prepared for Cynbit Technologies

Prepared by Surya Prakash Polkam • 20 October 2026

Assessment deliverable 01 | Confidential submission draft

AI Opportunity Report

Industry context, company alignment and recommended use case

Prepared for Cynbit Technologies • Prepared by Surya Prakash Polkam • 20 October 2026

Assessment deliverable 01 | Confidential submission draft

Executive summary

This report evaluates an AI-assisted project-inquiry intake workflow as a potential opportunity aligned with Cynbit's publicly described work in AI automation, LLM automation, CRM automation, custom software and cloud systems. The proposed workflow structures incoming inquiries, identifies missing details, drafts a follow-up and recommends a human-reviewed next action. LeadIQ is a proof of concept for that workflow.

Evidence boundary: public materials show Cynbit's services and a project inquiry touchpoint; they do not prove an internal lead-triage problem, quantify manual effort, or establish that the existing process is deficient. The opportunity is therefore a hypothesis to validate with stakeholders.

1. Company and workflow context

Cynbit's public website describes AI agents, LLM automation, CRM automation, custom software and cloud services. Its public contact page provides a project inquiry path. These make inquiry intake a relevant area to explore, without assuming that the company lacks existing capabilities.

- Public service alignment: LLM automation and AI-enabled workflows.
- Public service alignment: CRM automation and lead-related workflow support.
- Public workflow touchpoint: prospects can submit a project inquiry for follow-up.
- Scope: support human discovery; do not replace human judgment or automatically reject leads.

2. Opportunity hypothesis

Project inquiries can vary in completeness. A reviewer may need to identify service category, requirements, budget, timeline, urgency and unanswered questions before deciding on a next step. LeadIQ tests whether AI-assisted extraction and deterministic checks can make that first review more consistent.

3. Candidate use cases and controls

Use case	Potential value	Control / risk	Priority
Inquiry structuring	Reduce repeated copying from free text into fields.	Keep original message; reviewer verifies extracted fields.	High
Missing-information prompts	Prepare discovery questions consistently.	Rules may not fit all cases; permit override.	High
Follow-up draft	Accelerate first-response drafting.	Draft-only; human reviews and edits before sending.	High
Readiness label	Provide a consistent triage cue.	Current classifier is weak and synthetic; never use as conversion score.	Experimental
CRM integration	Reduce duplicate entry and support handoffs.	Requires authentication, audit, schema mapping and integration testing.	Later phase

4. Proposed workflow

– Receive the original inquiry text.

– Use an LLM to extract structured fields into a validated schema.

– Apply rules to flag missing budget, timeline, requirements or service category.

– Generate an experimental readiness category and state its limitations.

– Show extraction, missing details, draft reply and suggested next action.

– Require human review before sending any external message.

– Store the analysis for audit and later evaluation.

5. Measurement plan

Metric	Definition	Measurement approach
Time to first triage	Elapsed time from inquiry receipt to structured review.	Compare current workflow baseline with assisted workflow.
Field extraction accuracy	Correct extracted fields divided by fields present in source.	Human-labelled representative sample.
Unsupported-field rate	Extracted claims not grounded in the inquiry.	Reviewer audit; separately count critical errors.
Clarification relevance	Whether requested details are actually missing.	Reviewer rubric for relevant, redundant or missed questions.
Draft acceptance/edit rate	How often a draft is accepted or materially changed.	Compare generated and approved drafts.
Human override rate	Frequency that reviewer changes suggested category/action.	Log reviewer decisions and reason codes.

6. Risks and governance

- Privacy: use synthetic examples in demos; do not submit real customer/confidential data to an unapproved provider or free tier.
- Hallucination: unknown fields should remain unknown; preserve original text and support reviewer comparison.
- Fairness: do not use protected or irrelevant personal attributes for lead prioritisation.
- Human control: no automatic customer messaging, pricing commitments or lead rejection in the PoC.
- Security: keep keys in environment variables, exclude .env from Git, apply access controls and retention rules.
- Model quality: current readiness classifier uses 16 synthetic examples and obtained 50% exploratory 3-fold CV accuracy; this is not production evidence.

7. Roadmap

Phase	Scope	Exit criteria
Prototype	LLM extraction, rules, baseline classifier, UI, API, local persistence.	Functional smoke tests; limits documented.
Shadow pilot	Approved, redacted or consented historical inquiries; no automatic decisions.	Reviewed labels and field-level evaluation.
Integration	CRM write-back, access controls, audit logs, monitoring and retries.	Security and integration review; rollback path.
Optimisation	Train/calibrate using representative reviewed data.	Predefined metrics, drift checks and owner sign-off.

Conclusion

LeadIQ is a plausible opportunity to validate because it connects Cynbit's public LLM/CRM automation positioning with a visible project-inquiry touchpoint. The next step is stakeholder discovery and measurement—not a claim that an internal problem is confirmed or that the current classifier is production-ready.

Sources

– Cynbit official website — https://www.cynbit.in/

– Cynbit services — https://www.cynbit.in/services

– Cynbit contact / inquiry page — https://www.cynbit.in/contact

– Cynbit LinkedIn company page — https://in.linkedin.com/company/cynbit-in